

Your phone has an AI chip. Every night, it does nothing.

Modern phones carry a dedicated AI chip — an NPU — capable of roughly 40 trillion operations per second. From the moment you plug the phone in and go to sleep, that chip sits completely idle for eight hours.

NeuroMesh rents out those idle hours. Phone owners earn credits overnight. Developers get a big batch of AI work done by borrowing many phones at once, instead of renting an expensive server.

PROBLEM • THE PHONE OWNER

An expensive chip that earns nothing

Every flagship ships an NPU that sits dark from midnight to morning — charging, on Wi-Fi, screen off, cool. The most abundant idle AI chip on earth, and no way to make a single rupee from it.

PLAINLY

You already paid for the chip. It earns you nothing while you sleep.

PROBLEM • THE DEVELOPER

Simple bulk work, priced like heavy work

A developer processing two million documents does not need a top-end AI server. They need throughput. They rent a scarce, costly GPU anyway, because nothing cheaper will take the job.

PLAINLY

Like hiring a crane to carry a thousand grocery bags.

THE IDEA

One app, two sides of a market

Provider mode lends the phone's NPU under strict conditions and earns credits. Consumer mode submits a batch, watches it split across the fleet, and pays in those same credits.

PLAINLY

Lend your phone at night, spend what it earned during the day.

✓ Working prototype already built

6,474 lines of code

30 automated tests, all passing

Installable Android app, 43 MB

Networks like this already exist.

Every one of them rents computers. We rent phones.

NETWORK	WHAT IT RENTS OUT	WHO CAN SUPPLY IT	WHAT YOU NEED TO JOIN
Vast.ai	Desktop & server graphics cards	Crypto miners, small hosts	A spare ₹1.5 lakh GPU
Akash	Data-centre graphics cards	People who own server racks	Infrastructure and uptime
io.net	Pooled GPU clusters	Data centres, mining farms	To already be a supplier
Folding@home	Home computer power	Volunteers	Nothing — but nobody is paid
NeuroMesh	Phone AI chips	Anyone with a phone and a charger	To install an app and go to bed

WHY NOW

"Phones are too weak" stopped being true

That was correct when phone AI chips managed 5 trillion operations per second and overheated in ninety seconds. Today's flagships do 40–50 trillion. And the perfect conditions already happen in every home, every night, free.

PLAINLY

The technology caught up. Nobody updated the assumption.

OUR HONEST CLAIM

Not "cheaper than a server" — a closed loop

We do not pretend to beat data-centre pricing on raw cost. We offer what they cannot: earn credits while you sleep, spend them by day to run models your own phone is too small to run alone.

PLAINLY

The supplier and the customer are the same person.

SUPPLY IS THE EASY PART

The problem that kills these networks does not apply to us

Every compute marketplace struggles to find suppliers — they need people who own expensive hardware. Our supplier just needs a phone they already own.

PLAINLY

No new hardware has to exist anywhere for this to scale.

What we deliberately do not do, and why: phones can run AI models but cannot train them — the chip is built for one and physically not the other. And we only accept work that splits into fully independent pieces, so no phone ever waits on another. Both limits are designed in on purpose.

One job, cut into slices, handed out one at a time so no two phones can take the same slice.

NPU

The AI-only chip inside the phone.

SHARD

One small slice of a big job — 8 items out of 120.

LEASE

A 30-second reservation on a slice.

CREDITS

The in-app currency: earned, then spent.



TWO KINDS OF MEMORY

A fast temporary store holds live work. A permanent database holds accounts and money.

PLAINLY

A whiteboard for today, a ledger book for the year.

RESERVATIONS, NOT TRUST

Each slice is held for 30 seconds. A watchdog puts expired holds straight back in the queue.

PLAINLY

Quiet for 30 seconds? We assume the worst.

THE MONEY CAN'T DRIFT

Whole numbers only. Exactly one piece of code may touch a balance, with an audit beside it.

PLAINLY

One cashier, one till, one audit.

PRIVACY & SAFETY

Data is locked until a phone holds the slice. Someone else's job returns "not found".

PLAINLY

"Not allowed" would confirm it exists.

We did not just design this. We built it, and then we tried to break it.

Everything below is real output from the system running today — not a plan, not a mock-up.

A FULL JOB, START TO FINISH

15 / 15

120 items cut into 15 slices, finished by 3 phones in 1.9 seconds. All 120 answers came back in the original order and the credit accounts balanced exactly.

WE YANKED A PHONE MID-JOB

8 / 8

A phone was pulled away holding slice #0. Its hold expired, the slice went back, another phone finished it. Nothing lost — and the phone that walked away was paid zero.

16 PHONES FOUGHT OVER 1 SLICE

16 → 1

Exactly one won. Two phones on one slice means paying twice for one piece of work, so we made that impossible rather than unlikely.

WEAK PHONES GET PROTECTED

0

Slices offered to an older phone from a job it could not handle. Each phone is judged on what it actually measures, never on its brand name.

WHAT THE PERSON WHO SUBMITTED THE JOB SEES



finished

running now

dropped, retrying

waiting

One square per slice, updating live on the phone screen as the fleet works. This is the demo: the judges watch squares turn green in real time.

BUILT AND VERIFIED TODAY

- The coordinator — ~20 web endpoints, live progress feed, automatic database setup.
- The scheduler — five scripts covering take, hold, finish, fail and recover.
- The Android app — one app, both modes, background service, real UI. 43 MB.
- Three real AI models bundled in, each fingerprint-checked before it may run.
- Test tooling — a fake fleet, including a switch that kills a phone mid-job.

What we have NOT proven yet, said plainly: the Android app compiles and installs, but we have never run it on an actual iQOO 15 — we do not own one. Getting real inference running on the real chip is hour 0–3 of the hackathon, and it is exactly why we want to be in the room.

The phone is not where we show the project. The phone is the project.

WHICH CHIP RUNS THE MODEL

Try the AI chip, then the graphics chip, then the processor

A phone has three chips that can run a model, best to worst. We ask for Qualcomm's AI chip first, fall back to the graphics chip, then the main processor. If one fails to start we quietly step down instead of crashing.

PLAINLY

A slow demo beats a dead demo. Something always runs.

THE MODELS WE SHIP

Three real AI models, shrunk to fit a phone

— Text understanding — turns a sentence into comparable numbers. 188 KB.

— Image recognition — real trained classifier, 1,000 categories. 3.9 MB.

— Settings search — scores AI configurations to find the best one.

PLAINLY

We tested two versions of the photo model and kept the one that did not break on older phones.

RESPECTING THE OWNER

The phone only works when it can genuinely spare the effort

Charging AND on free Wi-Fi AND screen off AND not warm. All four, re-checked before every single slice, on the phone and again on the server. We stop before the phone gets hot enough to notice — not after.

PLAINLY

Pick up your phone and it stops instantly. Your battery, data and evening are never touched.

WHY THIS IS GENUINELY PHONE-FIRST

- The phone is the engine — take the phones away and there is no product.
- The phone is even the server. No laptop: a script runs the whole coordinator on one iQOO 15.
- The phone is the control panel — submitting, watching, earnings, all on-device.
- The phone set the rules. The 30-second hold and heat limit were chosen for phones.

WHAT WE WILL MEASURE AT THE HACKATHON

- Real numbers from the real chip — same slice on AI chip vs graphics chip vs processor, with speed and battery cost published.
- An hour-long heat test to tune our temperature limit against how the iQOO 15 really behaves.
- A big-screen judging view: mirror the live progress grid using vivo's multi-screen feature.

Most teams will spend the first day building the hard middle. Ours is already done.

HOURS	THE MILESTONE	WHERE WE ALREADY STAND	SO WE SPEND THE TIME ON
0-3	Get an AI model running on the phone itself	Not tried on real hardware	First real run on the iQOO 15. Our only unproven step.
3-8	Background service that fetches and runs work	Written, not device-tested	Prove the backup path on hardware, then the AI chip.
8-14	Splitting, scheduling, reassembling, paying	BUILT AND TESTED	Six hours we get back — measuring the chip, tuning for heat.
14-20	End to end: one phone, then many	Proven with fake phones	Swap the fake fleet for four real ones. Nothing else changes.
20-26	Job-submission screen and live progress	Built	Polish, the big-screen judging view, earnings history.
26-30	Rehearse, break it on purpose, keep a buffer	Break-test automated	Yank a phone on stage, on purpose, and let it recover live.

CAN IT GROW?

Designed for a crowd, not for four phones

Handing out a slice costs the same whether ten phones are asking or ten thousand. The server never tracks who is alive — the hold clock does it free. Going from 4 phones to 4,000 means a bigger queue and more copies of the same server.

PLAINLY

Nothing in the design gets slower as more people join.

WHO WOULD ACTUALLY USE IT

Both sides are paid in the same currency

A student processes a research dataset overnight using credits their own phone earned while charging. A phone owner turns eight dead hours into something spendable. Where renting a GPU costs real money and a phone is already in every pocket, that swap matters.

TEAM ← REPLACE THESE

Who is building it

- [Name] — Android & on-device AI: the app, the chips, the background service.
- [Name] — Server & distributed systems: scheduling, holds, credits.
- [Name] — Product & interface: the screens, the live view, the demo.
- [Name] — Infrastructure & testing: databases, break-tests, rehearsal.